

```
import pandas as pd
from pandas import DataFrame
import gspread
from google.oauth2.service_account import Credentials
from google.oauth2 import service_account
from pandas_gbq import to_gbq
```

▼ Important Functions

For production use only (example)

```
def get_google_sheet_as_df(sheet_id: str, sheet_name: str, creds_path: str) -> pd.DataFrame:
    """
    Retrieve data from a Google Sheet and return it as a Pandas DataFrame.

    :param sheet_id: The Google Sheet ID (found in the URL after /d/ and before /edit)
    :param sheet_name: The name of the sheet (tab) inside the document
    :param creds_path: Path to your service account credentials JSON file
    :return: Pandas DataFrame with the sheet data
    """
    scopes = ["https://www.googleapis.com/auth/spreadsheets.readonly"]
    credentials = Credentials.from_service_account_file(creds_path, scopes=scopes)
    client = gspread.authorize(credentials)

    sheet = client.open_by_key(sheet_id).worksheet(sheet_name)
    data = sheet.get_all_records()

    return pd.DataFrame(data)

def load_to_bigquery(df: DataFrame, table_name: str, credentials_path: str, if_exists: str = 'replace'):
    """
    Uploads a DataFrame to Google BigQuery using a service account credential file.

    Args:
        df (DataFrame): The pandas DataFrame to upload.
        table_name (str): Table name in the format 'project_id.dataset.table'.
        credentials_path (str): Path to the service account JSON credentials file.
        if_exists (str): What to do if the table exists: 'fail', 'replace', or 'append'. Default is 'replace'.

    Raises:
        ValueError: If the DataFrame is empty.
        Exception: If the upload fails.
    """
    if df.empty:
        raise ValueError("The DataFrame is empty. Nothing will be uploaded to BigQuery.")

    try:
        credentials = service_account.Credentials.from_service_account_file(credentials_path)
        to_gbq(df, table_name, credentials=credentials, if_exists=if_exists)
        print(f"✅ Data successfully uploaded to {table_name}.")
    except Exception as e:
        print(f"❌ Failed to upload data to BigQuery: {e}")
```

▼ Extraction

📁 Data Setup Instructions

The script expects a `data/` folder located in the **root of the GitHub repository**:

🔗 <https://github.com/JoaoPGOS/leanlayer-study-case.git>

This folder should contain the following CSV files:

- deals_meta.csv
- deals_snapshot.csv
- owners.csv
- targets.csv

✅ Structure Example

```
leanlayer-study-case/
├── data/
```

```

| | deals_meta.csv
| | deals_snapshot.csv
| | owners.csv
| | targets.csv
| etl.py

```

```

# In a production environment, it should use the function get_google_sheet_as_df using a google service account to extract this data
deals_meta_df = pd.read_csv('data/deals_meta.csv')
deals_snapshot_df = pd.read_csv('data/deals_snapshot.csv')
owners_df = pd.read_csv('data/owners.csv')
targets_df = pd.read_csv('data/targets.csv')

```

```
print(f"Log type: info\n\nDeals Meta\n{deals_meta_df.head()}\nDeals Snapshot\n{deals_snapshot_df.head()}\nOwners\n{owners_df.head()}\nTargets\n{targets_df.head()}")
```



Log type: info

Deals Meta

	deal_id	owner_id	created_date	close_date	deal_stage	forecast_category	\
0	deal_31	owner_1	2024-12-21	2025-06-02	Closed Lost		Closed
1	deal_34	owner_1	2024-05-29	2024-11-11	Prospecting		Pipeline
2	deal_65	owner_1	2024-08-26	2025-02-04	Prospecting		Pipeline
3	deal_78	owner_1	2024-02-21	2024-03-03	Closed Won		Closed
4	deal_92	owner_1	2024-06-08	2024-11-24	Qualification		Pipeline

	record_type	deal_type	deal_order_type	account_industry	\
0	Channel	Existing Business	Upsell		Tech
1	Direct	Existing Business	Cross-Sell		Retail
2	Direct	Existing Business	Cross-Sell		Healthcare
3	Channel	New Business	New Client		Finance
4	Channel	New Business	New Client		Healthcare

	account_region	account_size	deal_source	deal_amount
0	LATAM	SMB	Outbound	22127
1	EMEA	SMB	Partner	23224
2	LATAM	SMB	Outbound	17680
3	EMEA	SMB	Outbound	53629
4	LATAM	SMB	Inbound	23492

Deals Snapshot

	deal_id	snapshot_date	stage	forecast_category	amount	close_date	\
0	deal_31	2024-12-23	Prospecting	Pipeline	25164	2025-07-02	
1	deal_31	2024-12-30	Prospecting	Pipeline	17950	2025-04-02	
2	deal_31	2025-01-06	Prospecting	Pipeline	24985	2025-05-02	
3	deal_31	2025-01-13	Prospecting	Pipeline	18770	2025-07-02	
4	deal_31	2025-01-20	Prospecting	Pipeline	25485	2025-08-02	

owner_id

0	owner_1
1	owner_1
2	owner_1
3	owner_1
4	owner_1

Owners

	owner_id	name	email	role_start_date	\
0	owner_1	Christopher Moore	danielledunn@gmail.com	2023-02-12	
1	owner_2	Amber Jenkins	afoster@yahoo.com	2023-06-21	
2	owner_3	Mark Lopez	halleric@yahoo.com	2023-09-22	
3	owner_4	Kristen Mann	mhunt@james.org	2022-12-26	
4	owner_5	Dawn Alexander	kelliparker@coleman.com	2023-07-27	

	role_end_date	role	manager	segment
0	NaN	Account Executive	Traci Price	SMB
1	NaN	Account Executive	Victoria Brown	SMB
2	NaN	Account Executive	Stacey Mason	SMB
3	NaN	Account Executive	Teresa Vaughn	SMB
4	NaN	Account Executive	Laura Sanchez	SMB

Targets

	quarter	owner_id	segment	target_amount
0	2024 Q1	owner_1	SMB	42000
1	2024 Q2	owner_1	SMB	48300
2	2024 Q3	owner_1	SMB	55545
3	2024 Q4	owner_1	SMB	63877
4	2025 Q1	owner_1	SMB	73458

✓ Static analysis

✓ Breakdown of Lost Deals in Q1 2025

This code processes deals that were **Closed Lost** in **Q1 2025** and analyzes their progression:

1. Parse Close Dates

- Ensures `close_date` is in datetime format for filtering.

2. Filter Lost Deals in Q1

- Defines the Q1 2025 range (2025-01-01 to 2025-03-31)
- Filters deals where deal_stage == "Closed Lost" and close_date is within Q1
- Extracts the unique deal_id s of these lost deals

3. Filter Snapshots for Those Deals

- Filters the deals_snapshot_df to include only rows related to the lost deals
- Converts snapshot_date to datetime and sorts snapshots chronologically by deal_id and snapshot_date

4. Extract First and Last Snapshot Info

- Groups by deal_id and retrieves:
 - First and last **stage**
 - First and last **forecast category**
 - Final **owner ID**
 - Final **deal amount**

5. Output

- Displays a DataFrame summarizing the lifecycle of each lost deal, useful for analyzing pipeline flow and potential misclassification.

```
#>>>>>>>> Filter deals Closed Lost in Q1 <<<<<<<<<<
deals_meta_df['close_date'] = pd.to_datetime(deals_meta_df['close_date'])

# Define Q1 range (adjust year if needed)
start_q1 = "2025-01-01"
end_q1 = "2025-03-31"

lost_deals_q1 = deals_meta_df[
    (deals_meta_df['close_date'].between(start_q1, end_q1)) &
    (deals_meta_df['deal_stage'] == "Closed Lost")
][['deal_id']].unique()

print(f"Total lost deals in Q1: {len(lost_deals_q1)}")

#>>>>>>>> Filter snapshots for those deals <<<<<<<<<<
lost_snapshots = deals_snapshot_df[deals_snapshot_df['deal_id'].isin(lost_deals_q1)].copy()

lost_snapshots['snapshot_date'] = pd.to_datetime(lost_snapshots['snapshot_date'])
lost_snapshots = lost_snapshots.sort_values(by=['deal_id', 'snapshot_date'])

# >>>>>>>> Get first & last snapshot for each deal <<<<<<<<<<
lost_deals_df = (
    lost_snapshots.groupby('deal_id')
    .agg(
        old_stage=('stage', 'first'),
        last_stage=('stage', 'last'),
        old_forecast_category=('forecast_category', 'first'),
        last_forecast_category=('forecast_category', 'last'),
        owner_id=('owner_id', 'last'),
        amount=('amount', 'last')
    )
    .reset_index()
)

# ---- 4) Result ----
print(lost_deals_df)
```

```
➡ Total lost deals in Q1: 17
```

	deal_id	old_stage	last_stage	old_forecast_category	\
0	deal_114	Prospecting	Closed Lost	Pipeline	
1	deal_147	Prospecting	Closed Lost	Pipeline	
2	deal_271	Prospecting	Closed Lost	Pipeline	
3	deal_296	Prospecting	Closed Lost	Pipeline	
4	deal_299	Prospecting	Closed Lost	Pipeline	
5	deal_314	Prospecting	Closed Lost	Pipeline	
6	deal_401	Prospecting	Closed Lost	Pipeline	
7	deal_410	Prospecting	Closed Lost	Pipeline	
8	deal_47	Prospecting	Closed Lost	Pipeline	
9	deal_487	Prospecting	Closed Lost	Pipeline	
10	deal_514	Prospecting	Closed Lost	Pipeline	
11	deal_556	Prospecting	Closed Lost	Pipeline	
12	deal_575	Prospecting	Closed Lost	Pipeline	
13	deal_604	Prospecting	Closed Lost	Pipeline	
14	deal_674	Prospecting	Closed Lost	Pipeline	
15	deal_676	Prospecting	Closed Lost	Pipeline	
16	deal_96	Prospecting	Closed Lost	Pipeline	

	last_forecast_category	owner_id	amount
0	Closed	owner_11	239748
1	Closed	owner_1	30694
2	Closed	owner_7	58571
3	Closed	owner_13	54979
4	Closed	owner_8	257671
5	Closed	owner_3	31818
6	Closed	owner_9	48684
7	Closed	owner_7	27056
8	Closed	owner_5	14266
9	Closed	owner_9	59781
10	Closed	owner_3	29119
11	Closed	owner_5	19764
12	Closed	owner_3	18281
13	Closed	owner_9	43355
14	Closed	owner_5	36002
15	Closed	owner_3	25931
16	Closed	owner_5	38931

✓ Analysis of Deals Closed as Lost in Q1 2025

This script identifies and analyzes deals that were marked as **"Closed Lost"** during **Q1 2025**:

1. Filter by Close Date & Stage

- Select deals with `close_date` between **2025-01-01** and **2025-03-31**
- Only include deals where `deal_stage` is **"Closed Lost"**

2. Snapshot Filtering

- Filter `deals_snapshot_df` to include only rows where `deal_id` is in the list of lost deals

3. Chronological Ordering

- Sort the snapshots by `deal_id` and `snapshot_date` to ensure proper sequence

4. First vs. Last Snapshot Comparison

- For each deal, extract:
 - `old_stage` and `last_stage`: to track stage movement
 - `old_forecast_category` and `last_forecast_category`: to assess forecasting consistency
 - Final `owner_id` and `amount`: for owner and deal size context

This analysis helps understand the evolution of lost deals and identify potential issues such as forecast misclassification or unexpected deal movement.

```
def compute_stage_metrics(deals_snapshot_df: pd.DataFrame, stall_days_threshold: int = 30):
```

```
    """
```

```
    Given a full deals snapshot dataframe, adds columns:
```

- 'stage_change': True when stage differs from previous for same deal
- 'stage_start_date': first date stage started (forward filled)
- 'days_in_stage': days spent in current stage
- 'stalled': True if `days_in_stage` >= `threshold`

```
    Returns the full dataframe with these columns.
```

```
    """
```

```
    df = deals_snapshot_df.copy()
```

```
    df['snapshot_date'] = pd.to_datetime(df['snapshot_date'])
```

```
    df = df.sort_values(by=['deal_id', 'snapshot_date'])
```

```
    # Detect stage changes
```

```
    df['stage_change'] = df.groupby('deal_id')['stage'].transform(lambda x: x != x.shift())
```

```
    # Compute stage start date (ffill where stage_change is True)
```

```
    df['stage_start_date'] = df.groupby('deal_id')['snapshot_date'].transform(
        lambda dates: dates.where(df['stage_change']).ffill()
    )
```

```
    # Calculate days in current stage
```

```
    df['days_in_stage'] = (df['snapshot_date'] - df['stage_start_date']).dt.days
```

```
    # Mark stalled deals
```

```
    df['stalled'] = df['days_in_stage'] >= stall_days_threshold
```

```
    return df
```

```
def get_stalled_result_df(deals_snapshot_df: pd.DataFrame, stall_days_threshold: int = 30):
```

```
    """
```

```
    Returns a reduced dataframe with only latest snapshot per deal,
```

```
    filtered by deals whose snapshot_date is in the current year.
```

```
    Includes only selected columns plus stalled flag.
```

```

"""
df = compute_stage_metrics(deals_snapshot_df, stall_days_threshold)

# Get current year
today = pd.Timestamp.today()
current_year = today.year

# Get latest snapshot per deal
latest_snapshots = df.sort_values(by='snapshot_date').groupby('deal_id').tail(1)

# Filter to only deals in current year based on snapshot_date
filtered = latest_snapshots[
    (latest_snapshots['snapshot_date'].dt.year == current_year)
]

# Select only relevant columns for AI analysis
columns = ['deal_id', 'stage', 'days_in_stage', 'owner_id', 'forecast_category', 'amount', 'stalled']
return filtered[columns]

# full_deals_df has all snapshots + metrics
full_deals_df = compute_stage_metrics(deals_snapshot_df, stall_days_threshold=30)

# stalled_result_df filtered by current quarter, latest snapshot per deal (for AI input)
stalled_result_df = get_stalled_result_df(deals_snapshot_df, stall_days_threshold=30)

print(stalled_result_df.head())
full_deals_df["lost_deals_q1"] = full_deals_df["deal_id"].isin(stalled_result_df["deal_id"]).map({True: "Yes", False: "No"})

print(full_deals_df.loc[
    (full_deals_df["lost_deals_q1"] == "Yes")
])

```

```

↩

```

	deal_id	stage	days_in_stage	owner_id	forecast_category	\
4541	deal_374	Prospecting	126	owner_4	Pipeline	
6894	deal_358	Closed Won	0	owner_7	Closed	
6671	deal_113	Proposal	0	owner_7	Best Case	
1704	deal_583	Negotiation	0	owner_11	Commit	
1741	deal_601	Qualification	0	owner_11	Pipeline	

	amount	stalled
4541	21761	True
6894	53559	False
6671	23953	False
1704	47192	False
1741	135126	False

	deal_id	snapshot_date	stage	forecast_category	amount	\
1318	deal_114	2024-10-21	Prospecting	Pipeline	239748	
1319	deal_114	2024-10-28	Prospecting	Pipeline	239748	
1320	deal_114	2024-11-04	Prospecting	Pipeline	239748	
1321	deal_114	2024-11-11	Qualification	Pipeline	239748	
1322	deal_114	2024-11-18	Qualification	Pipeline	239748	
...	
5050	deal_96	2025-02-17	Proposal	Best Case	38931	
5051	deal_96	2025-02-24	Proposal	Best Case	38931	
5052	deal_96	2025-03-03	Negotiation	Best Case	38931	
5053	deal_96	2025-03-10	Closed Lost	Closed	38931	
5054	deal_96	2025-03-17	Closed Lost	Closed	38931	

	close_date	owner_id	stage_change	stage_start_date	days_in_stage	\
1318	2025-02-28	owner_11	True	2024-10-21	0	
1319	2025-02-28	owner_11	False	2024-10-21	7	
1320	2024-11-30	owner_11	False	2024-10-21	14	
1321	2024-11-30	owner_11	True	2024-11-11	0	
1322	2024-12-30	owner_11	False	2024-11-11	7	
...	
5050	2025-04-18	owner_5	True	2025-02-17	0	
5051	2025-01-18	owner_5	False	2025-02-17	7	
5052	2025-02-18	owner_5	True	2025-03-03	0	
5053	2025-01-18	owner_5	True	2025-03-10	0	
5054	2025-03-18	owner_5	False	2025-03-10	7	

	stalled	lost_deals_q1
1318	False	Yes
1319	False	Yes
1320	False	Yes
1321	False	Yes
1322	False	Yes
...
5050	False	Yes
5051	False	Yes
5052	False	Yes
5053	False	Yes
5054	False	Yes

[244 rows x 12 columns]

✓ Static Analysis

✓ Build Pacing Table for Next Quarter

This function estimates sales performance per owner for the **next quarter**, based on current pipeline status and static probability weights.

1. Active Owners

- Filters `owners_df` to include only those whose role is active or ends **after** the start of the next quarter.

2. Prepare Deals

- Converts `snapshot_date` to `datetime`.
- Reclassifies 'Discovery' stage to 'Prospecting'.
- Filters deals with snapshots in the **current year**.

3. Completion Rate

- Calculates completion rate per owner = $\text{Closed Won deals} \div \text{total deals}$ (clipped between 0 and 1).

4. Static Weights

- Defines **static probabilities** for each stage and forecast category:
 - `stage_prob`: e.g., 'Proposal' = 0.7, 'Closed Won' = 1.0
 - `forecast_weight`: e.g., 'Commit' = 0.9, 'Pipeline' = 0.3

5. Filter Valid Open Deals

- Removes deals in any 'Closed' stage.
- Excludes stalled deals (those stuck in a stage longer than `stall_days_threshold`).

6. Apply Weights & Estimate Value

- Gets the **latest snapshot** for each open deal.
- Applies:
 - `stage_prob`
 - `forecast_weight`
 - `owner_completion_rate`
- Computes $\text{adjusted_prob} = \text{stage_prob} * \text{forecast_weight} * \text{completion_rate}$
- Estimates $\text{value} = \text{amount} * \text{adjusted_prob}$

7. Aggregate by Owner

- Sums estimated pipeline value and open deal count per owner.

8. Merge with Targets

- Merges owner data with target data for the **next quarter**.
- Computes:
 - $\text{pacing} = \text{estimated_value} / \text{target_amount}$
 - $\text{gap} = \text{estimated_value} - \text{target_amount}$

9. Predictive Classification

- Classifies whether the owner is likely to hit their target using simple rules:
 - "Yes": if $\text{estimated} \geq \text{target}$
 - "Probably Yes": if $\text{pacing} \geq 0.5$ and has open deals
 - "Probably No": if $\text{pacing} < 0.5$ and has open deals
 - "No": no open deals

Returns:

A table with:

- Owner ID, Name, Segment
- Target and Estimated Values
- Pacing, Gap
- Open Deals Count
- Classification of likelihood to hit the target

```

import pandas as pd
from pandas.tseries.offsets import QuarterEnd

def build_pacing_table(owners_df, full_deals_df, targets_df, stall_days_threshold=30):
    today = pd.Timestamp.today()
    next_quarter = (today.month - 1) // 3 + 2
    next_year = today.year + (1 if next_quarter == 5 else 0)
    if next_quarter == 5:
        next_quarter = 1

    # --- 1) Active owners ---
    owners_active = owners_df[
        (owners_df['role_end_date'].isna()) |
        (pd.to_datetime(owners_df['role_end_date'], errors='coerce') >= pd.Timestamp(f"{next_year}-01-01"))
    ]

    # --- 2) Prepare deals ---
    df = full_deals_df.copy()
    df['snapshot_date'] = pd.to_datetime(df['snapshot_date'], errors='coerce')
    df['stage'] = df['stage'].replace({'Discovery': 'Prospecting'})
    current_year = today.year
    deals_this_year = df[df['snapshot_date'].dt.year == current_year]

    # --- 3) Owner completion rate ---
    total_deals = df.groupby("owner_id")["deal_id"].nunique()
    closed_deals = df[df['stage'].str.contains("Closed Won", case=False, na=False)] \
        .groupby("owner_id")["deal_id"].nunique()
    completion_rate = (closed_deals / total_deals).fillna(0).clip(0, 1)

    # --- 4) STATIC stage & forecast weights ---
    stage_prob = {
        'Proposal': 0.7,
        'Negotiation': 0.8,
        'Contract Sent': 0.9,
        'Qualification': 0.4,
        'Prospecting': 0.3,
        'Closed Won': 1.0,
        'Closed Lost': 0.0
    }
    forecast_weight = {
        'Commit': 0.9,
        'Best Case': 0.7,
        'Pipeline': 0.3,
        'Closed': 1.0
    }

    # --- 5) Filter valid open deals ---
    open_deals = deals_this_year[
        ~deals_this_year['stage'].str.contains("Closed", case=False, na=False)
    ]
    filtered_deals = open_deals[
        ~(open_deals['stalled'] == True) & (open_deals['days_in_stage'] > stall_days_threshold)
    ]

    # Latest snapshot per deal
    latest_deals = (
        filtered_deals
        .sort_values('snapshot_date')
        .groupby('deal_id')
        .tail(1)
        .set_index('deal_id')
    )

    # --- 6) Apply static probabilities ---
    latest_deals['stage_prob'] = latest_deals['stage'].map(stage_prob).fillna(0.2)
    latest_deals['forecast_weight'] = latest_deals['forecast_category'].map(forecast_weight).fillna(0.1)
    latest_deals['completion_rate'] = latest_deals['owner_id'].map(completion_rate).fillna(0)

    latest_deals['adjusted_prob'] = (
        latest_deals['stage_prob'] * latest_deals['forecast_weight'] * latest_deals['completion_rate']
    )
    latest_deals['estimated_value'] = latest_deals['amount'] * latest_deals['adjusted_prob']

    # --- 7) Aggregate by owner ---
    expected_pipeline = latest_deals.groupby('owner_id')['estimated_value'].sum().reset_index()
    open_deals_count = latest_deals.groupby('owner_id').size().reset_index(name='open_deals_count')

    # Targets for the next quarter
    target_quarter_str = f"{next_year} Q{next_quarter}"
    targets_next_q = targets_df[targets_df['quarter'] == target_quarter_str]

    # --- 8) Merge everything ---
    pacing_df = (

```

```

    owners_active
    .merge(targets_next_q[['owner_id', 'target_amount']], on='owner_id', how='left')
    .merge(expected_pipeline, on='owner_id', how='left')
    .merge(open_deals_count, on='owner_id', how='left')
)

pacing_df['estimated_value'] = pacing_df['estimated_value'].fillna(0)
pacing_df['open_deals_count'] = pacing_df['open_deals_count'].fillna(0).astype(int)
pacing_df['pacing'] = pacing_df['estimated_value'] / pacing_df['target_amount']
pacing_df['gap'] = pacing_df['estimated_value'] - pacing_df['target_amount']

# --- 9) Simplified classification (Yes, Probably Yes, Probably No, No) ---
end_of_quarter = today + QuarterEnd(0)
days_remaining_in_quarter = (end_of_quarter - today).days

def classify_hit(row):
    if row["estimated_value"] >= row["target_amount"]:
        return "Yes"
    # Heuristic based only on pacing & deal count
    if row["pacing"] >= 0.5 and row["open_deals_count"] > 0:
        return "Probably Yes"
    if row["pacing"] < 0.5 and row["open_deals_count"] > 0:
        return "Probably No"
    return "No"

pacing_df["will_hit_target"] = pacing_df.apply(classify_hit, axis=1)

return pacing_df[['owner_id', 'name', 'segment', 'target_amount',
                  'estimated_value', 'pacing', 'gap',
                  'open_deals_count', 'will_hit_target']]

pacing_df = build_pacing_table(owners_df, full_deals_df, targets_df)
pacing_df["pacing"] = pacing_df["pacing"].round(2).astype(float)

```

pacing_df

	owner_id	name	segment	target_amount	estimated_value	pacing	gap	open_deals_count	will_hit_target
0	owner_1	Christopher Moore	SMB	111721	30266.732075	0.27	-81454.267925	25	Probably No
1	owner_2	Amber Jenkins	SMB	111721	79201.428980	0.71	-32519.571020	24	Probably Yes
2	owner_3	Mark Lopez	SMB	111721	47120.025455	0.42	-64600.974545	16	Probably No
3	owner_4	Kristen Mann	SMB	111721	54769.281538	0.49	-56951.718462	22	Probably No
4	owner_5	Dawn Alexander	SMB	111721	92643.665068	0.83	-19077.334932	34	Probably Yes
5	owner_6	Kelly Lin	Mid-Market	138321	109736.456923	0.79	-28584.543077	21	Probably Yes
6	owner_7	Dr. Kimberly Brennan	Mid-Market	138321	86818.062857	0.63	-51502.937143	19	Probably Yes
7	owner_8	Allen Gilmore	Mid-Market	138321	70248.582326	0.51	-68072.417674	18	Probably Yes
8	owner_9	Barbara Robinson	Mid-Market	138321	87188.497200	0.63	-51132.502800	25	Probably Yes
9	owner_10	Anthony Jackson	Enterprise	156941	31806.725385	0.20	-125134.274615	14	Probably No
10	owner_11	Michael Williams	Enterprise	156941	55410.930789	0.35	-101530.069211	13	Probably No

✓ Clean

✓ Final Pacing Data Preparation & Formatting

This section finalizes the `pacing_df` and adds helpful formatting and ordering for deal stage progression analysis.

1. Merge Owner Names into Full Deals

- Merges `pacing_df[['owner_id', 'name']]` into `full_deals_df` to include owner names.

2. Stage Ordering for Analysis

- Defines a custom sort order for deal stages using the `stage_order` dictionary:
 - e.g., `'Prospecting' = 'a'`, `'Closed Lost' = 'f'`
- Adds a new column `order` to `full_deals_df` based on this mapping.

3. Round Numeric Columns

- Rounds numeric columns (`estimated_value`, `gap`, `target_amount`, `pacing`) in `pacing_df` to **2 decimal places** and casts them as `float`.

4. Format Columns for Display

- Formats the main numeric columns to **two decimal places with comma as decimal separator** (e.g., `1.23` → `"1,23"`) for user-friendly display:
 - Columns: `"pacing"`, `"target_amount"`, `"estimated_value"`, `"gap"`

This ensures the final pacing table is both accurate and presentation-ready.

```
full_deals_df = full_deals_df.merge(
    pacing_df[["owner_id", "name"]],
    on="owner_id",
    how="left"
)

stage_order = {
    "Prospecting": "a",
    "Qualification": "b",
    "Proposal": "c",
    "Negotiation": "d",
    "Closed Won": "e",
    "Closed Lost": "f"
}

full_deals_df["order"] = full_deals_df["stage"].map(stage_order)

pacing_df[["estimated_value", "gap", "target_amount", "pacing"]] = pacing_df[["estimated_value", "gap", "target_amount", "pacing"]].round(2)

cols_to_format = ["pacing", "target_amount", "estimated_value", "gap"]
for col in cols_to_format:
    pacing_df[col] = pacing_df[col].apply(lambda x: f"{x:.2f}".replace(".", ",") if isinstance(x, float) else x)
pacing_df
```

	owner_id	name	segment	target_amount	estimated_value	pacing	gap	open_deals_count	will_hit_target
0	owner_1	Christopher Moore	SMB	111721,00	30266,73	0,27	-81454,27	25	Probably No
1	owner_2	Amber Jenkins	SMB	111721,00	79201,43	0,71	-32519,57	24	Probably Yes
2	owner_3	Mark Lopez	SMB	111721,00	47120,03	0,42	-64600,97	16	Probably No
3	owner_4	Kristen Mann	SMB	111721,00	54769,28	0,49	-56951,72	22	Probably No
4	owner_5	Dawn Alexander	SMB	111721,00	92643,67	0,83	-19077,33	34	Probably Yes
5	owner_6	Kelly Lin	Mid-Market	138321,00	109736,46	0,79	-28584,54	21	Probably Yes
6	owner_7	Dr. Kimberly Brennan	Mid-Market	138321,00	86818,06	0,63	-51502,94	19	Probably Yes
7	owner_8	Allen Gilmore	Mid-Market	138321,00	70248,58	0,51	-68072,42	18	Probably Yes
8	owner_9	Barbara Robinson	Mid-Market	138321,00	87188,50	0,63	-51132,50	25	Probably Yes
9	owner_10	Anthony Jackson	Enterprise	156941,00	31806,73	0,20	-125134,27	14	Probably No
10	owner_11	Michael Williams	Enterprise	156941,00	55410,93	0,35	-101530,07	13	Probably No
11	owner_12	Amanda Hutchinson	Strategic	170241,00	146908,22	0,86	-23332,78	24	Probably Yes
12	owner_13	Eileen Nunez	Strategic	170241,00	180273,04	1,06	10032,04	24	Yes

▼ Load

```
full_deals_df.to_csv("full_deals.csv", index=False)
pacing_df.to_csv("pacing.csv", index=False)
```

